

L3 Exercise 3 - Parallel ETL - Solution

June 7, 2023

1 Exercise 3: Parallel ETL

```
In [1]: %load_ext sql
In [2]: from time import time
import configparser
import matplotlib.pyplot as plt
import pandas as pd
```

2 STEP 1: Get the params of the created redshift cluster

- We need:
 - The redshift cluster endpoint
 - The IAM role ARN that give access to Redshift to read from S3

```
In [3]: config = configparser.ConfigParser()
config.read_file(open('dwh.cfg'))
KEY=config.get('AWS','key')
SECRET= config.get('AWS','secret')

DWH_DB= config.get("DWH","DWH_DB")
DWH_DB_USER= config.get("DWH","DWH_DB_USER")
DWH_DB_PASSWORD= config.get("DWH","DWH_DB_PASSWORD")
DWH_PORT = config.get("DWH","DWH_PORT")

In [4]: # FILL IN THE REDSHIFT ENPOINT HERE
# e.g. DWH_ENDPOINT="redshift-cluster-1.csmamz5zxmle.us-west-2.redshift.amazonaws.com"
DWH_ENDPOINT="dwhcluster.c4uipqmqcj1l.us-west-2.redshift.amazonaws.com"

#FILL IN THE IAM ROLE ARN you got in step 2.2 of the previous exercise
#e.g DWH_ROLE_ARN="arn:aws:iam::988332130976:role/dwhRole"
DWH_ROLE_ARN="arn:aws:iam::918744264023:role/dwhRole"
```

3 STEP 2: Connect to the Redshift Cluster

```
In [5]: conn_string="postgresql://{user}:{password}@{host}/{db}".format(DWH_DB_USER, DWH_DB_PASSWORD, DWH_ENDPOINT)
print(conn_string)
%sql $conn_string
```

```
postgresql://dwhuser:PasswOrd@dwhcluster.c4uipqmqcj1l.us-west-2.redshift.amazonaws.com:5439/dwh
```

```
Out[5]: 'Connected: dwhuser@dwh'
```

```
In [6]: import boto3
```

```
s3 = boto3.resource('s3',  
                    region_name="us-west-2",  
                    aws_access_key_id=KEY,  
                    aws_secret_access_key=SECRET  
                    )
```

```
sampleDbBucket = s3.Bucket("udacity-labs")
```

```
for obj in sampleDbBucket.objects.filter(Prefix="tickets"):  
    print(obj)
```

```
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/full/')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/full/full.csv.gz')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00000-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00001-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00002-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00003-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00004-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00005-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00006-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00007-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00008-d33afb94-b8af-407d-ab')  
s3.ObjectSummary(bucket_name='udacity-labs', key='tickets/split/part-00009-d33afb94-b8af-407d-ab')
```

4 STEP 3: Create Tables

```
In [7]: %%sql
```

```
DROP TABLE IF EXISTS "sporting_event_ticket";  
CREATE TABLE "sporting_event_ticket" (  
    "id" double precision DEFAULT nextval('sporting_event_ticket_seq') NOT NULL,  
    "sporting_event_id" double precision NOT NULL,  
    "sport_location_id" double precision NOT NULL,  
    "seat_level" numeric(1,0) NOT NULL,  
    "seat_section" character varying(15) NOT NULL,  
    "seat_row" character varying(10) NOT NULL,  
    "seat" character varying(10) NOT NULL,  
    "ticketholder_id" double precision,  
    "ticket_price" numeric(8,2) NOT NULL  
);
```

```
* postgresql://dwhuser:***@dwhcluster.c4uipqmqcj1l.us-west-2.redshift.amazonaws.com:5439/dwh
Done.
Done.
```

```
Out[7]: []
```

5 STEP 4: Load Partitioned data into the cluster

```
In [8]: %%time
        qry = """
            copy sporting_event_ticket from 's3://udacity-labs/tickets/split/part'
            credentials 'aws_iam_role={}'
            gzip delimiter ';' compupdate off region 'us-west-2';
        """.format(DWH_ROLE_ARN)

        %sql $qry

* postgresql://dwhuser:***@dwhcluster.c4uipqmqcj1l.us-west-2.redshift.amazonaws.com:5439/dwh
Done.
CPU times: user 0 ns, sys: 0 ns, total: 0 ns
Wall time: 11.7 s
```

6 STEP 4: Create Tables for the non-partitioned data

```
In [9]: %%sql
        DROP TABLE IF EXISTS "sporting_event_ticket_full";
        CREATE TABLE "sporting_event_ticket_full" (
            "id" double precision DEFAULT nextval('sporting_event_ticket_seq') NOT NULL,
            "sporting_event_id" double precision NOT NULL,
            "sport_location_id" double precision NOT NULL,
            "seat_level" numeric(1,0) NOT NULL,
            "seat_section" character varying(15) NOT NULL,
            "seat_row" character varying(10) NOT NULL,
            "seat" character varying(10) NOT NULL,
            "ticketholder_id" double precision,
            "ticket_price" numeric(8,2) NOT NULL
        );

* postgresql://dwhuser:***@dwhcluster.c4uipqmqcj1l.us-west-2.redshift.amazonaws.com:5439/dwh
Done.
Done.
```

```
Out[9]: []
```

7 STEP 5: Load non-partitioned data into the cluster

- Note how it's slower than loading partitioned data

```
In [10]: %%time
```

```
qry = """
    copy sporting_event_ticket_full from 's3://udacity-labs/tickets/full/full.csv.gz'
    credentials 'aws_iam_role={}'
    gzip delimiter ';' compupdate off region 'us-west-2';
    """.format(DWH_ROLE_ARN)
```

```
%sql $qry
```

```
* postgresql://dwhuser:***@dwhcluster.c4uipqmqcj1l.us-west-2.redshift.amazonaws.com:5439/dwh
Done.
```

```
CPU times: user 0 ns, sys: 0 ns, total: 0 ns
```

```
Wall time: 22.5 s
```

```
In [ ]:
```